

```

1  """
2  train_model.py
3  =====
4  Description:
5      Defines the training configuration, default hyperparameters, and execution
6      entry point for the Neural Sequence Decoder. Implements command-line
7      argument parsing for hyperparameter sweeps (ablation studies) and
8      initiates the training process.
9
10     Part of the final project for ECE C143A/C243A at UCLA.
11
12 Authors:
13     Yihe Xu, Injun Choi, Saiya Shah
14
15 Date:
16     December 1, 2025
17
18 Location:
19     neural_seq_decoder/scripts/train_model.py
20 """
21
22
23 import sys
24
25 modelName = 'speechBaseline0'
26
27 args = {}
28 args['outputDir'] = '../src_bk/speech_logs/' + modelName
29 args['datasetPath'] = '../src_bk/ptDecoder_ctc.pkl'
30 args['seqLen'] = 150
31 args['maxTimeSeriesLen'] = 1200
32 args['batchSize'] = 64
33 args['lrStart'] = 0.02
34 args['lrEnd'] = 0.02
35 args['nUnits'] = 1024
36 args['nBatch'] = 10000
37 args['nLayers'] = 5
38 args['seed'] = 0
39 args['nClasses'] = 40
40 args['nInputFeatures'] = 256
41 args['dropout'] = 0.4
42 args['whiteNoiseSD'] = 0.8
43 args['constantOffsetSD'] = 0.2
44 args['gaussianSmoothWidth'] = 2.0
45 args['strideLen'] = 4
46 args['kernelLen'] = 32
47 args['bidirectional'] = False
48 args['l2_decay'] = 1e-5
49
50 # ---- NEW AUGMENTATION SETTINGS ----
51 args['timeMask'] = -1        # maximum mask length
52 args['timeMask_N'] = -1      # number of masks
53 args['channelDrop'] = 0.4    # drop 20% of channels
54
55
56 argv = sys.argv
57
58 def get_arg(flag, cast_type, default):
59     """Return parsed CLI argument or fallback to default."""
60     if flag in argv:
61         idx = argv.index(flag)
62         if idx + 1 < len(argv):
63             return cast_type(argv[idx + 1])
64     return default
65

```

```
66 # map the flags your sweep sends into your args dictionary
67 args['whiteNoiseSD'] = get_arg("--white_std", float, args['whiteNoiseSD'])
68 args['constantOffsetSD'] = get_arg("--mean_std", float, args['constantOffsetSD'])
69 args['gaussianSmoothWidth'] = get_arg("--gauss_sigma", float, args['gaussianSmoothWidth'])
70 args['kernelLen'] = get_arg("--gauss_kernel", int, args['kernelLen'])
71 args['timeMask'] = get_arg("--mask_max_len", int, args['timeMask'])
72 args['timeMask_N'] = get_arg("--mask_n", int, args['timeMask_N']) # NEW
73 args['channelDrop'] = get_arg("--drop_p", float, args['channelDrop'])

74
75
76 from neural_decoder.neural_decoder_trainer import trainModel
77
78 trainModel(args)
```

```

1  """
2  augmentation.py
3  =====
4  Description:
5      Defines PyTorch modules for data augmentation and signal processing
6      specifically tailored for neural sequence decoding.
7
8  Includes:
9      - WhiteNoise: Additive Gaussian noise.
10     - MeanDriftNoise: Channel-wise constant offset noise.
11     - GaussianSmoothing: 1D/2D smoothing via convolution.
12     - TimeMask: SpecAugment-style temporal masking.
13     - ChannelDropout: Randomly drops input channels (simulating electrode failure).
14
15     Part of the final project for ECE C143A/C243A at UCLA.
16
17 Authors:
18     Yihe Xu, Injun Choi, Saiya Shah
19
20 Date:
21     December 1, 2025
22
23 Location:
24     neural_seq_decoder/src/neural_decoder/augmentation.py
25 """
26
27
28 import math
29 import numbers
30 import torch
31 from torch import nn
32 from torch.nn import functional as F
33
34
35 class WhiteNoise(nn.Module):
36     def __init__(self, std=0.1):
37         super().__init__()
38         self.std = std
39
40     def forward(self, x):
41         noise = torch.randn_like(x) * self.std
42         return x + noise
43
44 class MeanDriftNoise(nn.Module):
45     def __init__(self, std=0.1):
46         super().__init__()
47         self.std = std
48
49     def forward(self, x):
50         _, C = x.shape
51         noise = torch.randn(1, C) * self.std
52         return x + noise
53
54 class GaussianSmoothing(nn.Module):
55     """
56     Apply gaussian smoothing on a
57     1d, 2d or 3d tensor. Filtering is performed separately for each channel
58     in the input using a depthwise convolution.
59     Arguments:
60         channels (int, sequence): Number of channels of the input tensors. Output will
61             have this number of channels as well.
62         kernel_size (int, sequence): Size of the gaussian kernel.
63         sigma (float, sequence): Standard deviation of the gaussian kernel.
64         dim (int, optional): The number of dimensions of the data.
65             Default value is 2 (spatial).

```

```

66     """
67
68     def __init__(self, channels, kernel_size, sigma, dim=2):
69         super(GaussianSmoothing, self).__init__()
70         if isinstance(kernel_size, numbers.Number):
71             kernel_size = [kernel_size] * dim
72         if isinstance(sigma, numbers.Number):
73             sigma = [sigma] * dim
74
75         # The gaussian kernel is the product of the
76         # gaussian function of each dimension.
77         kernel = 1
78         meshgrids = torch.meshgrid(
79             [torch.arange(size, dtype=torch.float32) for size in kernel_size]
80         )
81         for size, std, mgrid in zip(kernel_size, sigma, meshgrids):
82             mean = (size - 1) / 2
83             kernel *= (
84                 1
85                 / (std * math.sqrt(2 * math.pi))
86                 * torch.exp(-(((mgrid - mean) / std) ** 2) / 2)
87             )
88
89         # Make sure sum of values in gaussian kernel equals 1.
90         kernel = kernel / torch.sum(kernel)
91
92         # Reshape to depthwise convolutional weight
93         kernel = kernel.view(1, 1, *kernel.size())
94         kernel = kernel.repeat(channels, *[1] * (kernel.dim() - 1))
95
96         self.register_buffer("weight", kernel)
97         self.groups = channels
98
99         if dim == 1:
100             self.conv = F.conv1d
101         elif dim == 2:
102             self.conv = F.conv2d
103         elif dim == 3:
104             self.conv = F.conv3d
105         else:
106             raise RuntimeError(
107                 "Only 1, 2 and 3 dimensions are supported. Received {}".format(dim)
108             )
109
110     def forward(self, input):
111         """
112             Apply gaussian filter to input.
113             Arguments:
114                 input (torch.Tensor): Input to apply gaussian filter on.
115             Returns:
116                 filtered (torch.Tensor): Filtered output.
117             """
118             return self.conv(input, weight=self.weight, groups=self.groups, padding="same")
119
120
121 class TimeMask(nn.Module):
122     """
123         SpecAugment-style time masking for neural time series.
124         Randomly masks out consecutive time steps.
125
126         Expected input shape: (C, T)
127         - C = channels/electrodes
128         - T = time steps
129     """
130     def __init__(self, max_mask_len=25, n_masks=2, mask_value=0.0):
131         super().__init__()
132         self.max_mask_len = max_mask_len

```

```

133         self.n_masks = n_masks
134         self.mask_value = mask_value
135
136     def forward(self, x):
137         """
138             Args:
139                 x: Tensor (C, T)
140             Returns:
141                 Masked tensor of same shape
142             """
143         if not self.training:
144             return x
145
146         C, T = x.shape
147         x = x.clone() # avoid in-place ops; safer for CTC & autograd
148
149         for _ in range(self.n_masks):
150             # random mask length
151             mask_len = torch.randint(
152                 1, self.max_mask_len + 1,
153                 (1,), device=x.device
154             ).item()
155
156             # random start index
157             max_start = max(1, T - mask_len)
158             start = torch.randint(
159                 0, max_start,
160                 (1,), device=x.device
161             ).item()
162
163             # apply mask
164             x[:, start:start + mask_len] = self.mask_value
165
166         return x
167
168
169     class ChannelDropout(nn.Module):
170         """
171             Randomly drops entire channels (electrodes).
172             Good for robustness to impedance spikes or bad contacts.
173             """
174         def __init__(self, p=0.1):
175             super().__init__()
176             self.p = p
177
178         def forward(self, x):
179             C = x.shape[0]
180             drop_mask = (torch.rand(C) > self.p).float().to(x.device)
181             drop_mask = drop_mask.unsqueeze(-1)
182             return x * drop_mask
183

```

```

1  """
2  dataset.py
3  =====
4  Description:
5      Defines the PyTorch Dataset class for loading neural activity and
6      aligned phoneme sequences. Handles multi-day recording sessions
7      by flattening the trial structure while preserving session (day)
8      indices for domain adaptation.
9
10     Part of the final project for ECE C143A/C243A at UCLA.
11
12 Authors:
13     Yihe Xu, Injun Choi, Saiya Shah
14
15 Date:
16     December 1, 2025
17
18 Location:
19     neural_seq_decoder/src/neural_decoder/dataset.py
20 """
21
22
23 import torch
24 from torch.utils.data import Dataset
25
26
27 class SpeechDataset(Dataset):
28     def __init__(self, data, transform=None):
29         self.data = data
30         self.transform = transform
31         self.n_days = len(data)
32         self.n_trials = sum([len(d["sentenceDat"]) for d in data])
33
34         self.neural_feats = []
35         self.phone_seqs = []
36         self.neural_time_bins = []
37         self.phone_seq_lens = []
38         self.days = []
39         for day in range(self.n_days):
40             for trial in range(len(data[day]["sentenceDat"])):
41                 self.neural_feats.append(data[day]["sentenceDat"][trial])
42                 self.phone_seqs.append(data[day]["phonemes"][trial])
43                 self.neural_time_bins.append(data[day]["sentenceDat"][trial].shape[0])
44                 self.phone_seq_lens.append(data[day]["phoneLens"][trial])
45                 self.days.append(day)
46
47     def __len__(self):
48         return self.n_trials
49
50     def __getitem__(self, idx):
51         neural_feats = torch.tensor(self.neural_feats[idx], dtype=torch.float32)
52
53         if self.transform:
54             neural_feats = self.transform(neural_feats)
55
56         return (
57             neural_feats,
58             torch.tensor(self.phone_seqs[idx], dtype=torch.int32),
59             torch.tensor(self.neural_time_bins[idx], dtype=torch.int32),
60             torch.tensor(self.phone_seq_lens[idx], dtype=torch.int32),
61             torch.tensor(self.days[idx], dtype=torch.int64),
62         )
63

```

```

1  """
2   model.py
3  =====
4  Description:
5      Defines the model architecture for the Neural Sequence Decoder.
6      Part of the final project for ECE C143A/C243A at UCLA.
7
8  Authors:
9      Yihe Xu, Injun Choi, Saiya Shah
10
11 Date:
12     December 1, 2025
13
14 Location:
15     neural_seq_decoder/src/neural_decoder/model.py
16 """
17
18
19 import torch
20 from torch import nn
21 from .augmentations import GaussianSmoothing
22
23 class GradReverse(torch.autograd.Function):
24     @staticmethod
25     def forward(ctx, x, lambd):
26         ctx.lambd = lambd
27         return x.view_as(x)
28
29     @staticmethod
30     def backward(ctx, grad_output):
31         return -ctx.lambd * grad_output, None
32
33 def grad_reverse(x, lambd=1.0):
34     return GradReverse.apply(x, lambd)
35
36 class GRUDecoder(nn.Module):
37     def __init__(
38         self,
39         neural_dim,
40         n_classes,
41         hidden_dim,
42         layer_dim,
43         nDays=24,
44         dropout=0,
45         device="cuda",
46         strideLen=4,
47         kernelLen=14,
48         gaussianSmoothWidth=0,
49         bidirectional=False,
50     ):
51         super(GRUDecoder, self).__init__()
52
53         # Define day classifier
54         feat_dim = hidden_dim * (2 if bidirectional else 1)
55         self.day_classifier = nn.Sequential(
56             nn.Linear(feat_dim, hidden_dim // 2),
57             nn.ReLU(),
58             nn.Linear(hidden_dim // 2, nDays),
59         )
60
61         # Defining the number of layers and the nodes in each layer
62         self.layer_dim = layer_dim
63         self.hidden_dim = hidden_dim
64         self.neural_dim = neural_dim
65         self.n_classes = n_classes

```

```

66         self.nDays = nDays
67         self.device = device
68         self.dropout = dropout
69         self.strideLen = strideLen
70         self.kernelLen = kernelLen
71         self.gaussianSmoothWidth = gaussianSmoothWidth
72         self.bidirectional = bidirectional
73         self.inputLayerNonlinearity = torch.nn.Softsign()
74         self.unfolder = torch.nn.Unfold(
75             (self.kernelLen, 1), dilation=1, padding=0, stride=self.strideLen
76         )
77         self.gaussianSmoothen = GaussianSmoothing(
78             neural_dim, 20, self.gaussianSmoothWidth, dim=1
79         )
80         self.dayWeights = torch.nn.Parameter(torch.randn(nDays, neural_dim, neural_dim))
81         self.dayBias = torch.nn.Parameter(torch.zeros(nDays, 1, neural_dim))
82
83     for x in range(nDays):
84         self.dayWeights.data[x, :, :] = torch.eye(neural_dim)
85
86     # GRU layers
87     self.gru_decoder = nn.GRU(
88         (neural_dim) * self.kernelLen,
89         hidden_dim,
90         layer_dim,
91         batch_first=True,
92         dropout=self.dropout,
93         bidirectional=self.bidirectional,
94     )
95
96     for name, param in self.gru_decoder.named_parameters():
97         if "weight_hh" in name:
98             nn.init.orthogonal_(param)
99         if "weight_ih" in name:
100             nn.init.xavier_uniform_(param)
101
102     # Input layers
103     for x in range(nDays):
104         setattr(self, "inpLayer" + str(x), nn.Linear(neural_dim, neural_dim))
105
106     for x in range(nDays):
107         thisLayer = getattr(self, "inpLayer" + str(x))
108         thisLayer.weight = torch.nn.Parameter(
109             thisLayer.weight + torch.eye(neural_dim)
110         )
111
112     # rnn outputs
113     if self.bidirectional:
114         self.fc_decoder_out = nn.Linear(
115             hidden_dim * 2, n_classes + 1
116         ) # +1 for CTC blank
117     else:
118         self.fc_decoder_out = nn.Linear(hidden_dim, n_classes + 1) # +1 for CTC blank
119
120     def forward(self, neuralInput, dayIdx):
121         neuralInput = torch.permute(neuralInput, (0, 2, 1))
122         neuralInput = self.gaussianSmoothen(neuralInput)
123         neuralInput = torch.permute(neuralInput, (0, 2, 1))
124
125         # apply day layer
126         dayWeights = torch.index_select(self.dayWeights, 0, dayIdx)
127         transformedNeural = torch.einsum(
128             "btd,bdk->btk", neuralInput, dayWeights
129         ) + torch.index_select(self.dayBias, 0, dayIdx)
130         transformedNeural = self.inputLayerNonlinearity(transformedNeural)
131
132         # stride/kernel

```

```
133     stridedInputs = torch.permute(
134         self.unfolder(
135             torch.unsqueeze(torch.permute(transformedNeural, (0, 2, 1)), 3)
136         ),
137         (0, 2, 1),
138     )
139
140     # apply RNN layer
141     if self.bidirectional:
142         h0 = torch.zeros(
143             self.layer_dim * 2,
144             transformedNeural.size(0),
145             self.hidden_dim,
146             device=self.device,
147         ).requires_grad_()
148     else:
149         h0 = torch.zeros(
150             self.layer_dim,
151             transformedNeural.size(0),
152             self.hidden_dim,
153             device=self.device,
154         ).requires_grad_()
155
156     hid, _ = self.gru_decoder(stridedInputs, h0.detach())
157
158     # get seq
159     seq_out = self.fc_decoder_out(hid)
160     rev = grad_reverse(hid, lambd=1.0)
161     day_logits = self.day_classifier(rev)
162     return seq_out, day_logits, hid
163     # return seq_out
164
```

```

1  """
2  neural_decoder_trainer.py
3  =====
4  Description:
5      Implements the core training loop, validation logic, and data loading
6      utilities for the Neural Sequence Decoder.
7
8      Key features include:
9      1. SWATS Optimization: Automates the switch from Adam to SGD based on
10         projected learning rates to improve generalization.
11      2. Domain Adversarial Training: Integrates gradient reversal for
12         day/domain independence.
13      3. Real-time Evaluation: Monitors Character Error Rate (CER) and
14         Connectionist Temporal Classification (CTC) loss.
15
16      Part of the final project for ECE C143A/C243A at UCLA.
17
18  Authors:
19      Yihe Xu, Injun Choi, Saiya Shah
20
21  Date:
22      December 1, 2025
23
24  Location:
25      neural_seq_decoder/src/neural_decoder/neural_decoder_trainer.py
26  """
27
28
29 import os
30 import pickle
31 import time
32
33 from edit_distance import SequenceMatcher
34 import hydra
35 import numpy as np
36 import torch
37 from torch.nn.utils.rnn import pad_sequence
38 from torch.utils.data import DataLoader
39 import torch.nn.functional as F
40
41 from .model import GRUDecoder, grad_reverse
42 from .dataset import SpeechDataset
43
44 import torch.optim.lr_scheduler
45
46 # --- PATCH START: Fix for PyTorch < 2.0 ---
47 if not hasattr(torch.optim.lr_scheduler, "LRScheduler"):
48     # Create the alias 'LRScheduler' pointing to the hidden '_LRScheduler'
49     torch.optim.lr_scheduler.LRScheduler = torch.optim.lr_scheduler._LRScheduler
50 # --- PATCH END ---
51
52 # Now you can safely import from the library
53 from pytorch_optimizer import Lookahead
54
55 def get_smart_sgd_lr(adam_optimizer, beta2=0.999):
56     """
57         Calculates the global SGD learning rate by projecting the
58         adaptive learning rates of Adam onto a single scalar.
59     """
60     print("[SWATS] Calculating projection from Adam state...")
61
62     # 1. Get the current step (needed for bias correction)
63     step = 0
64     for group in adam_optimizer.param_groups:
65         for p in group['params']:

```

```

66         if p.grad is None: continue
67         state = adam_optimizer.state[p]
68         if 'step' in state:
69             step = state['step']
70             if isinstance(step, torch.Tensor):
71                 step = step.item()
72             break
73         if step > 0: break
74
75     if step == 0:
76         print("[SWATS] Warning: Optimizer has no step count. Defaulting to 0.01")
77         return 0.01
78
79     # 2. Calculate Bias Correction for the Second Moment (v_t)
80     bias_correction = 1 - beta2 ** step
81
82     effective_lrs = []
83
84     # 3. Iterate over all parameters to find mean effective step size
85     for group in adam_optimizer.param_groups:
86         base_lr = group['lr']
87         eps = group['eps'] # Use the epsilon defined in the optimizer
88
89         for p in group['params']:
90             if p.grad is None: continue
91
92             state = adam_optimizer.state[p]
93             if 'exp_avg_sq' not in state: continue
94
95             # Get the variance vector (v)
96             v = state['exp_avg_sq']
97
98             # Correct bias: v_hat = v / (1 - beta2^t)
99             v_hat = v / bias_correction
100
101            # Calculate Effective LR vector: base_lr / (sqrt(v_hat) + eps)
102            # We assume the step direction is roughly consistent
103            layer_effective_lr = base_lr / (torch.sqrt(v_hat) + eps)
104
105            # We take the mean of this tensor to get a scalar for this parameter
106            effective_lrs.append(layer_effective_lr.mean().item())
107
108    if len(effective_lrs) == 0:
109        return 0.01
110
111    # 4. Project to a single scalar (Arithmetic Mean of all effective LRs)
112    projected_lr = np.mean(effective_lrs)
113    return projected_lr
114
115
116 def getDatasetLoaders(
117     datasetName,
118     batchSize,
119 ):
120     with open(datasetName, "rb") as handle:
121         loadedData = pickle.load(handle)
122
123     def _padding(batch):
124         X, y, X_lens, y_lens, days = zip(*batch)
125         X_padded = pad_sequence(X, batch_first=True, padding_value=0)
126         y_padded = pad_sequence(y, batch_first=True, padding_value=0)
127
128         return (
129             X_padded,
130             y_padded,
131             torch.stack(X_lens),
132             torch.stack(y_lens),

```

```

133         torch.stack(days),
134     )
135
136     train_ds = SpeechDataset(loadedaData["train"], transform=None)
137     test_ds = SpeechDataset(loadedaData["test"])
138
139     train_loader = DataLoader(
140         train_ds,
141         batch_size=batchSize,
142         shuffle=True,
143         num_workers=0,
144         pin_memory=True,
145         collate_fn=_padding,
146     )
147     test_loader = DataLoader(
148         test_ds,
149         batch_size=batchSize,
150         shuffle=False,
151         num_workers=0,
152         pin_memory=True,
153         collate_fn=_padding,
154     )
155
156     return train_loader, test_loader, loadedaData
157
158 def trainModel(args):
159     os.makedirs(args["outputDir"], exist_ok=True)
160     torch.manual_seed(args["seed"])
161     np.random.seed(args["seed"])
162     device = "cuda"
163
164     with open(args["outputDir"] + "/args", "wb") as file:
165         pickle.dump(args, file)
166
167     trainLoader, testLoader, loadedaData = getDatasetLoaders(
168         args["datasetPath"],
169         args["batchSize"],
170     )
171
172     model = GRUDecoder(
173         neural_dim=args["nInputFeatures"],
174         n_classes=args["nClasses"],
175         hidden_dim=args["nUnits"],
176         layer_dim=args["nLayers"],
177         nDays=len(loadedaData["train"]),
178         dropout=args["dropout"],
179         device=device,
180         strideLen=args["strideLen"],
181         kernelLen=args["kernelLen"],
182         gaussianSmoothWidth=args["gaussianSmoothWidth"],
183         bidirectional=args["bidirectional"],
184     ).to(device)
185
186     loss_ctc = torch.nn.CTCLoss(blank=0, reduction="mean", zero_infinity=True)
187     optimizer = torch.optim.Adam(
188         model.parameters(),
189         lr=args["lrStart"],
190         betas=(0.9, 0.999),
191         eps=0.1,
192         weight_decay=args["l2_decay"],
193     )
194
195     # Learning Rate Scheduler: Linear Decay from lrStart to lrEnd
196     scheduler = torch.optim.lr_scheduler.LinearLR(
197         optimizer,
198         start_factor=1.0,
199         end_factor=args["lrEnd"] / args["lrStart"],

```

```

200         total_iters=args["nBatch"],
201     )
202
203     # --- SWATS CONFIGURATION ---
204     # Switch at 40% or 50% of batches usually works best for projection
205     SWITCH_BATCH = int(args["nBatch"] * 0.3)
206     switched_to_sgd = False
207
208     # --train--
209     testLoss = []
210     testCER = []
211     startTime = time.time()
212     for batch in range(args["nBatch"]):
213         if batch == SWITCH_BATCH and not switched_to_sgd:
214             print("-" * 50)
215             print(f"Batch {batch}: Initiating SWATS (Smart Switch to SGD)...")
216
217         # A. Calculate the projected Learning Rate
218         # We pass 0.999 as beta2 because that matches your Adam init above
219         smart_sgd_lr = get_smart_sgd_lr(optimizer, beta2=0.999)
220
221         print(f"Adam LR was: {args['lrStart']} ")
222         print(f"Projected SGD LR is: {smart_sgd_lr:.6f}")
223
224         # B. Re-initialize Optimizer as SGD
225         optimizer = torch.optim.SGD(
226             model.parameters(),
227             lr=smart_sgd_lr,
228             momentum=0.9, # Momentum is critical for SGD
229             weight_decay=args["l2_decay"],
230         )
231
232         # C. Re-initialize Scheduler
233         # We need a new scheduler that decays from the NEW lr to the FINAL lr
234         # over the remaining batches.
235         remaining_iters = args["nBatch"] - batch
236
237         # Calculate what the end factor should be relative to the new SGD start
238         # If we want to end at args["lrEnd"], the factor is lrEnd / smart_sgd_lr
239         new_end_factor = args["lrEnd"] / smart_sgd_lr
240
241         scheduler = torch.optim.lr_scheduler.LinearLR(
242             optimizer,
243             start_factor=1.0,
244             end_factor=new_end_factor,
245             total_iters=remaining_iters,
246         )
247
248         switched_to_sgd = True
249         print(f"Switched to SGD with LR={smart_sgd_lr:.5f}, Momentum=0.9")
250         print("-" * 50)
251
252         model.train()
253
254         X, y, X_len, y_len, dayIdx = next(iter(trainLoader))
255         X, y, X_len, y_len, dayIdx = (
256             X.to(device),
257             y.to(device),
258             X_len.to(device),
259             y_len.to(device),
260             dayIdx.to(device),
261         )
262
263         # Noise augmentation is faster on GPU
264         if args["whiteNoiseSD"] > 0:
265             X += torch.randn(X.shape, device=device) * args["whiteNoiseSD"]
266

```

```

267     if args["constantOffsetSD"] > 0:
268         X += (
269             torch.randn([X.shape[0], 1, X.shape[2]], device=device)
270             * args["constantOffsetSD"]
271         )
272
273     pred, day_logits, hid = model.forward(X, dayIdx)
274
275     ctc_loss = loss_ctc(
276         torch.permute(pred.log_softmax(2), [1, 0, 2]),
277         y,
278         ((X_len - model.kernelLen) / model.strideLen).to(torch.int32),
279         y_len,
280     )
281     ctc_loss = torch.sum(ctc_loss)
282
283     # Day classification loss (adversarial)
284     B, T, feat_dim = hid.shape
285     rev_features = grad_reverse(hid, lambd=1.0)
286     day_logits = model.day_classifier(rev_features.reshape(-1, feat_dim))
287     day_targets = dayIdx.unsqueeze(1).repeat(1, T).reshape(-1)
288     day_loss = F.cross_entropy(day_logits, day_targets)
289
290     # Combined loss
291     alpha = 0.1
292     loss = ctc_loss + alpha * day_loss
293
294     # Backpropagation
295     optimizer.zero_grad()
296     loss.backward()
297     optimizer.step()
298     scheduler.step()
299
300     # print(endTime - startTime)
301
302     # Eval
303     if batch % 100 == 0:
304         with torch.no_grad():
305             model.eval()
306             allLoss = []
307             total_edit_distance = 0
308             total_seq_length = 0
309             for X, y, X_len, y_len, testDayIdx in testLoader:
310                 X, y, X_len, y_len, testDayIdx = (
311                     X.to(device),
312                     y.to(device),
313                     X_len.to(device),
314                     y_len.to(device),
315                     testDayIdx.to(device),
316                 )
317
318                 # pred = model.forward(X, testDayIdx)
319                 pred, day_logits, hid = model.forward(X, testDayIdx)
320                 loss = loss_ctc(
321                     torch.permute(pred.log_softmax(2), [1, 0, 2]),
322                     y,
323                     ((X_len - model.kernelLen) / model.strideLen).to(torch.int32),
324                     y_len,
325                 )
326                 loss = torch.sum(loss)
327                 allLoss.append(loss.cpu().detach().numpy())
328
329                 adjustedLens = ((X_len - model.kernelLen) / model.strideLen).to(
330                     torch.int32
331                 )
332                 for iterIdx in range(pred.shape[0]):
333                     decodedSeq = torch.argmax(

```

```

334             torch.tensor(pred[iterIdx, 0 : adjustedLens[iterIdx], :]),
335             dim=-1,
336         ) # [num_seq]
337         decodedSeq = torch.unique_consecutive(decodedSeq, dim=-1)
338         decodedSeq = decodedSeq.cpu().detach().numpy()
339         decodedSeq = np.array([i for i in decodedSeq if i != 0])
340
341         trueSeq = np.array(
342             y[iterIdx][0 : y_len[iterIdx]].cpu().detach()
343         )
344
345         matcher = SequenceMatcher(
346             a=trueSeq.tolist(), b=decodedSeq.tolist()
347         )
348         total_edit_distance += matcher.distance()
349         total_seq_length += len(trueSeq)
350
351         avgDayLoss = np.sum(allLoss) / len(testLoader)
352         cer = total_edit_distance / total_seq_length
353
354         endTime = time.time()
355         print(
356             f"batch {batch}, ctc loss: {avgDayLoss:>7f}, cer: {cer:>7f}, time/batch: {((endTime
- startTime)/100):>7.3f}"
357         )
358         startTime = time.time()
359
360         if len(testCER) > 0 and cer < np.min(testCER):
361             torch.save(model.state_dict(), args["outputDir"] + "/modelWeights")
362             testLoss.append(avgDayLoss)
363             testCER.append(cer)
364
365         tStats = {}
366         tStats["testLoss"] = np.array(testLoss)
367         tStats["testCER"] = np.array(testCER)
368
369         with open(args["outputDir"] + "/trainingStats", "wb") as file:
370             pickle.dump(tStats, file)
371
372
373     def loadModel(modelDir, nInputLayers=24, device="cuda"):
374         modelWeightPath = modelDir + "/modelWeights"
375         with open(modelDir + "/args", "rb") as handle:
376             args = pickle.load(handle)
377
378         model = GRUDecoder(
379             neural_dim=args["nInputFeatures"],
380             n_classes=args["nClasses"],
381             hidden_dim=args["nUnits"],
382             layer_dim=args["nLayers"],
383             nDays=nInputLayers,
384             dropout=args["dropout"],
385             device=device,
386             strideLen=args["strideLen"],
387             kernelLen=args["kernelLen"],
388             gaussianSmoothWidth=args["gaussianSmoothWidth"],
389             bidirectional=args["bidirectional"],
390         ).to(device)
391
392         model.load_state_dict(torch.load(modelWeightPath, map_location=device))
393         return model
394
395
396     @hydra.main(version_base="1.1", config_path="conf", config_name="config")
397     def main(cfg):
398         cfg.outputDir = os.getcwd()
399         trainModel(cfg)

```

```
400
401 if __name__ == "__main__":
402     main()
```